

1. Introducing data analysis in R

Frank Edwards

9/2/2019

Introduction to Statistics

Why learn statistics?

1. Quantitative methods allow us to discover or infer patterns when we have large amounts of data
2. Statistics provide methods for testing for differences between groups of data
3. Always remember two things: 1) all models are wrong, but some are useful; 2) social data come from people, and are always imperfect

1. Introduce students to statistical computing through the R programming language
2. Develop data manipulation, exploration, and visualization skills
3. Introduce core concepts in probability and statistics

- Lecture introduces new materials and concepts
- Lab provides interactive time to work on programming
- Office hours (FE: Friday 9:30-11, RL: Monday 11 - 1)
- Raven's office hours provide a great opportunity to get help with the homework

Weekly routine for the course

1. Do the assigned readings
2. Attend lecture
3. Practice skills in lab
4. Do homework
5. Take a break
6. return to 1

1. Imai, *Quantitative Social Science*

This book is the foundation of the course. It introduces core social science methods, programming and statistics.

Homework

- I will assign a problem set at the end of each lecture
- Problem sets will be posted at github.com/f-edwards/intro_stats/
- Data is posted either on GitHub or available through the QSS package
- Problem sets are due at 10AM the day of the following lecture
- Explain your answers in writing. Don't just submit code.
- Homework is hard! Give it plenty of time.
- Email me (frank.edwards@rutgers.edu) the compiled RMarkdown output file and source code (.html and .rmd files) by the due date

- Course website:
`https://f-edwards.github.io/intro_stats/`
- Course Slack: `https://introstatsfall2020.slack.com`

Questions about the course?

- R: <https://cran.r-project.org/>
- RStudio: <https://www.rstudio.com/>

Packages for the course

```
#install.packages(c("devtools", "tidyverse"))  
library(devtools)  
### For book data  
install_github("kosukeimai/qss-package",  
               build_vignettes = TRUE)  
### For problem set data  
install_github("conjugateprior/qss.student")
```

Problems?

Introducing: R!

```
5+3
```

```
## [1] 8
```

R is a calculator

```
5*3
```

```
## [1] 15
```



```
5/3
```

```
## [1] 1.666667
```

```
5^3
```

```
## [1] 125
```

R can make comparisons

```
5>3
```

```
## [1] TRUE
```

R can make comparisons

```
5<=3
```

```
## [1] FALSE
```

R can make comparisons

```
5==3
```

```
## [1] FALSE
```

R can make comparisons

```
"a"=="a"
```

```
## [1] TRUE
```

R works with objects

```
a<-2
```

```
a+1
```

```
## [1] 3
```

R works with objects

```
a<-2
```

```
a+1
```

```
## [1] 3
```



```
a<-2
```

```
a+1
```

```
## [1] 3
```

```
b<-a+2
```

```
a<-2
```

```
a<-a+1
```

```
a
```

```
## [1] 3
```

Objects can take many types

```
a<-2
```

```
class(a)
```

```
## [1] "numeric"
```

```
b<-"howdy"
```

```
class(b)
```

```
## [1] "character"
```

```
c<-TRUE
```

```
class(c)
```

```
## [1] "logical"
```

Objects can take many types

```
b<-"howdy"
```

```
class(b)
```

```
## [1] "character"
```

Objects can take many types

```
c<-TRUE  
class(c)
```

```
## [1] "logical"
```

Vectors are one-dimensional arrays of values of any class

```
vector1<-c(2,3,4,5,6)
```

```
vector1
```

```
## [1] 2 3 4 5 6
```

Vectors are one-dimensional arrays of values of any class

```
vector2<-c("a", "fancy", "vector")
```

```
vector2
```

```
## [1] "a"      "fancy"  "vector"
```



```
vector3<-c(TRUE, FALSE, TRUE, FALSE, FALSE)  
vector3
```

```
## [1] TRUE FALSE TRUE FALSE FALSE
```

Vectorized operations

```
vector1
```

```
## [1] 2 3 4 5 6
```

```
2 * vector1
```

```
## [1] 4 6 8 10 12
```

```
vector3
```

```
## [1] TRUE FALSE TRUE FALSE FALSE
```

```
vector3==FALSE
```

```
## [1] FALSE TRUE FALSE TRUE TRUE
```

Vectorized operations

```
vector3
```

```
## [1] TRUE FALSE TRUE FALSE FALSE
```

```
vector3==FALSE
```

```
## [1] FALSE TRUE FALSE TRUE TRUE
```

```
vector2
```

```
## [1] "a"      "fancy"  "vector"
```

```
vector2
```

```
## [1] "a"      "fancy"  "vector"
```

```
vector2[1]
```

```
## [1] "a"
```

```
vector2
```

```
## [1] "a"      "fancy"  "vector"
```

```
vector2[2]
```

```
## [1] "fancy"
```

```
vector2
```

```
## [1] "a"      "fancy"  "vector"
```

```
vector2[3]
```

```
## [1] "vector"
```

```
vector1
```

```
## [1] 2 3 4 5 6
```

```
vector1[2] + 3
```

```
## [1] 6
```


Functions!

R has loads and loads of functions.

- Functions run a fixed set of operations on some argument(s)
- Functions return a value that can be assigned to an object
- Functions take the general form *function(arguments)*

Some common function in R for vectors

```
vector1
```

```
## [1] 2 3 4 5 6
```

```
min(vector1)
```

```
## [1] 2
```

```
max(vector1)
```

```
## [1] 6
```

```
mean(vector1)
```

```
## [1] 4
```

```
sum(vector1)
```

Some common function in R for vectors

```
vector1
```

```
## [1] 2 3 4 5 6
```

```
max(vector1)
```

```
## [1] 6
```

Some common function in R for vectors

```
vector1
```

```
## [1] 2 3 4 5 6
```

```
mean(vector1)
```

```
## [1] 4
```

Some common function in R for vectors

```
vector1
```

```
## [1] 2 3 4 5 6
```

```
sum(vector1)
```

```
## [1] 20
```

Functions can work together

```
sum(vector1)
```

```
## [1] 20
```

```
length(vector1)
```

```
## [1] 5
```

```
sum(vector1)/length(vector1)
```

```
## [1] 4
```

We can even write our own!

```
redundantMean<-function(x){  
  n<-length(x)  
  sum_x<-sum(x)  
  xbar<-sum_x/n  
  return(xbar)  
}  
redundantMean(vector1)
```

```
## [1] 4
```

Questions?

```
### load the book package
library(qss)
### attach the UNpop data
data(UNpop)
head(UNpop)
```

```
##   year world.pop
## 1 1950   2525779
## 2 1960   3026003
## 3 1970   3691173
## 4 1980   4449049
## 5 1990   5320817
## 6 2000   6127700
```

As super-matrices?

Recall that we can obtain any element x_{ij} from a matrix X with row index i and column index j

UNpop

##	year	world.pop
## 1	1950	2525779
## 2	1960	3026003
## 3	1970	3691173
## 4	1980	4449049
## 5	1990	5320817
## 6	2000	6127700
## 7	2010	6916183

As super-matrices?

Recall that we can obtain any element x_{ij} from a matrix X with row index i and column index j

```
UNpop[1,1]
```

```
## [1] 1950
```

As super-matrices?

Recall that we can obtain any element x_{ij} from a matrix X with row index i and column index j

```
UNpop[2,2]
```

```
## [1] 3026003
```

Other means of indexing

```
UNpop[,1]
```

```
## [1] 1950 1960 1970 1980 1990 2000 2010
```

Other means of indexing

```
UNpop[1,]
```

```
##   year world.pop
```

```
## 1 1950   2525779
```

Other means of indexing

```
UNpop$year
```

```
## [1] 1950 1960 1970 1980 1990 2000 2010
```

Other means of indexing

```
UNpop$year[3]
```

```
## [1] 1970
```


Other means of indexing

```
UNpop[1, "world.pop"]
```

```
## [1] 2525779
```

Some convenient data.frame functions

```
summary(UNpop)
```

##	year	world.pop
##	Min. :1950	Min. :2525779
##	1st Qu.:1965	1st Qu.:3358588
##	Median :1980	Median :4449049
##	Mean :1980	Mean :4579529
##	3rd Qu.:1995	3rd Qu.:5724258
##	Max. :2010	Max. :6916183

Some convenient data.frame functions

```
nrow(UNpop)
```

```
## [1] 7
```

```
ncol(UNpop)
```

```
## [1] 2
```

```
dim(UNpop)
```

```
## [1] 7 2
```

Some convenient data.frame functions

```
head(UNpop)
```

```
##   year world.pop  
## 1 1950   2525779  
## 2 1960   3026003  
## 3 1970   3691173  
## 4 1980   4449049  
## 5 1990   5320817  
## 6 2000   6127700
```

Some convenient data.frame functions

```
tail(UNpop)
```

```
##   year world.pop  
## 2 1960   3026003  
## 3 1970   3691173  
## 4 1980   4449049  
## 5 1990   5320817  
## 6 2000   6127700  
## 7 2010   6916183
```

Some convenient data.frame functions

```
names(UNpop)
```

```
## [1] "year"      "world.pop"
```

How to do operations over a data.frame column

```
mean(UNpop$world.pop)
```

```
## [1] 4579529
```

```
sum(UNpop$world.pop)/nrow(UNpop)
```

```
## [1] 4579529
```

How to do operations over a data.frame column

```
UNpop$world.pop/UNpop$world.pop[1]
```

```
## [1] 1.000000 1.198047 1.461400 1.761456 2.106604 2.426063
```


Homework 1

- Complete HW1, posted at https://github.com/f-edwards/intro_stats/tree/master/hw
- Due Wednesday, September 9 at 10AM
- Submit homework as .Rmd and .html to frank.edwards@rutgers.edu
- I encourage you to work in groups, but make sure you submit your own code and write-up
- #HW1 is open on Slack for Q&A. Use it!
- You can access the data through the QSS package (see how I loaded the UNpop data above)

Lab

- RMarkdown allows us to combine code and text in one document
- With the proper workflow, you can do all your academic writing in one place!
- Let's work through a demo!

```
install.packages("swirl")
```

```
library(swirl) # load the swirl package  
install_course_github("kosukeimai", "qss-swirl")
```

```
swirl()
```

```
library(swirl)  
swirl()
```