

## Measurement and visualization, 2

---

Frank Edwards

10/1/2019

## Visualizing data

- `ggplot`

```
ggplot(ipv, ## object, usually a data.frame)  
  aes(x = region) + ## aesthetic variables, generally x, y, color, etc  
  geom_bar() ## a geom to plot the aesthetics
```

```
## Note that + in ggplot() works the same way as %>% in tidyverse:  
## It strings together commands, evaluated in sequence
```

- `geom_bar()`, `geom_histogram()`, `geom_boxplot()`, `geom_point()`
- Wickham Chapter 3:  
<https://r4ds.had.co.nz/data-visualisation.html>
- Recommended reading: Kieran Healy, *Data Visualization*. Available free at socviz.co
- Available workshops on tidyverse, ggplot, rmarkdown at New Brunswick

## Today: Measurement in the social sciences

- Survey methods with randomization
- Administrative data and agency surveys
- Unit, item non-response
- Desirability bias
- Latent variables, latent groups
- More visualization

- A census records information about a population, with measurement for each individual or unit in the population
- A survey samples from a population to make an inference about population characteristics

## The basic motivation for survey sampling

```
n <- 1e+06
marbles <- data.frame(color = c(rep("blue", n), rep("green", n),
  n)))
table(marbles$color)


##
##      blue      green      red
## 1000000 1000000 1000000
```

How could we know how many of each color are in the enormous bag of marbles?

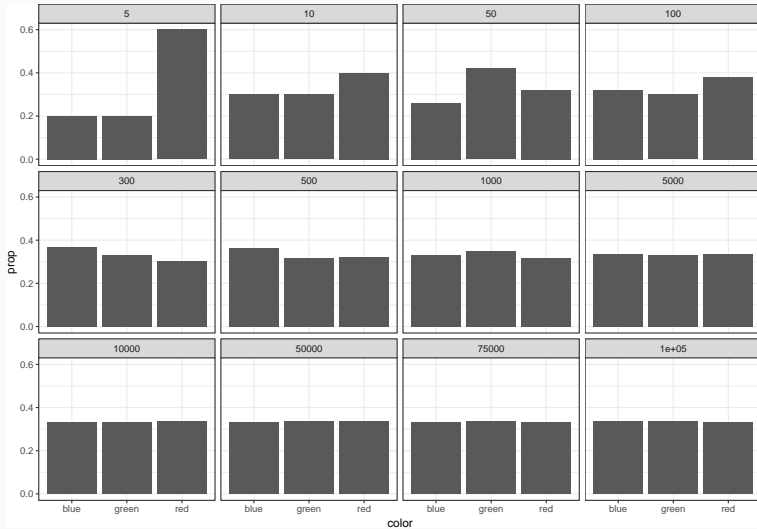
- Count them all (tedious!)
- Sample

# The Truth

```
ggplot(marbles, aes(x = color)) + geom_bar()
```



How many random draws is enough to accurately measure the characteristics of 3 million marbles?





With a sufficiently large sample and equal probability of sampling for all units in the population, a simple random sample allows for unbiased measurement of population characteristics.

Such a sample is representative of the population across both measured and unmeasured characteristics

If you are sampling, it should (in general) be randomized

## Stratified random sampling

If we wish to learn about particular sub-populations (i.e. geographies), we can use multi-stage or stratified sampling

1. Randomly sample larger units (geographic) or select larger units of interest purposively
2. Randomly sample individuals within these larger units

EXAMPLE: The American Community Survey (simplified)

1. Take a list of all US Census tracts
2. Randomly sample households within tract based on complete list of addresses (sampling frame)
3. Randomly sample adults within household, conduct survey

## When surveys go wrong

1. Unit non-response
2. Item non-response
3. Lying

Individual (or organization) doesn't respond to the survey

- How are surveys actually administered?
- Response rates are generally low (and decreasing!)

Completely random non-response is not a problem.

Individual (or organization) doesn't respond to the survey

- How are surveys actually administered?
- Response rates are generally low (and decreasing!)

Completely random non-response is not a problem.

When would non-response be an issue?

Individual takes the survey, but refuses to answer (skips) a particular question

- Why might this occur?

Individual takes the survey, but refuses to answer (skips) a particular question

- Why might this occur?
- When would this be a problem?

## Lying (ok... misrepresentation)

- Social desirability bias
  - Did you vote? Remember HW 1?
  - Are you a racist?
  - What kinds of crimes do you like to do?



Examining non-response in a survey  
of exposure to violence in  
Afghanistan

---

## Load the data

```
library(qss)  
data(afghan)  
data(afghan.village)
```

# Explore the variables in afghan

```
table(afghan$province)
```


```
##  
## Helmand   Khost   Kunar   Logar Uruzgan  
##      855     630     396     486     387
```

```
table(afghan$district)
```

```
##  
##      Asadabad      Bak Baraki Barak      Chapa Dara      Dangam  
##          54          54          180          108          63  
##      Dihrawud      Garmser      Ghaziabad Khas Uruzgan      Khoshi  
##          117          225          63          117          54  
##          Khost      Lashkar Gah      Musa Qala      Naw Zad      Puli Alam  
##          243          108          225          216          252  
##      Qalandar Shahidi Hassas      Spira      Tani      Washer  
##          63          153          99          171          81  
##      Wata Pur  
##          108
```


## Explore the variables in afghan

```
ggplot(afghan, aes(x = age)) + geom_histogram()
```



## Explore the variables in afghan

```
ggplot(afghan, aes(x = educ.years)) + geom_bar()
```



## Explore the variables in afghan

```
table(afghan$employed)
```

```
##
```

```
##      0      1
```

```
## 1149 1605
```

# Explore the variables in afghan

```
table(afghan$income)
```

```
##  
## 10,001-20,000 2,001-10,000 20,001-30,000 less than 2,000  
##          616          1420           93          457  
## over 30,000  
##          14
```

```
## for ordered categorical
```

```
afghan <- afghan %>% mutate(income = factor(income, levels = c("less than 2,000",  
  "2,001-10,000", "10,001-20,000", "20,001-30,000", "over 30,000")))
```

```
table(afghan$income)
```

```
##  
## less than 2,000 2,001-10,000 10,001-20,000 20,001-30,000  
##          457          1420           616           93  
## over 30,000  
##          14
```

## Explore the variables in afghan

```
afghan %>% select(employed, violent.exp.ISAF, violent.exp.taliban)
```

##	employed	violent.exp.ISAF	violent.exp.taliban
##	Min. :0.0000	Min. :0.0000	Min. :0.0000
##	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000
##	Median :1.0000	Median :0.0000	Median :0.0000
##	Mean :0.5828	Mean :0.3749	Mean :0.3289
##	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000
##	Max. :1.0000	Max. :1.0000	Max. :1.0000
##		NA's :25	NA's :54




## Explore the variables in afghan.village

```
head(afghan.village)
```

```
##      altitude population village.surveyed
## 1  1959.08         197              1
## 2  2425.88         744              0
## 3  2236.60         179              1
## 4  1691.76         225              0
## 5  1928.04         379              0
## 6  1194.56         617              0
```


## Explore the variables in afghan.village

```
ggplot(afghan.village, aes(x = altitude)) + geom_histogram()
```




## Explore the variables in afghan.village

```
ggplot(afghan.village, aes(x = population)) + geom_histogram()
```



## Explore the variables in afghan.village: logs help!

```
ggplot(afghan.village, aes(x = log(population))) + geom_histogram
```




## Explore the variables in afghan.village

```
mean(afghan.village$village.surveyed)
```

```
## [1] 0.1094421
```


## Is the sampling representative of villages?

```
ggplot(afghan.village, aes(x = village.surveyed == 1, y = altitude))
```




## Is the sampling representative of villages?

```
ggplot(afghan.village, aes(x = village.surveyed == 1, y = log(po  
  geom_boxplot())
```



# Does item non-response bias estimates of violence by region?





- Unit non-responses can bias survey estimates
- Item non-response can bias survey estimates
- Social desirability can bias survey estimates
- Errors induced by these biases can lead to incorrect conclusions (see polling consensus on 2016 election)

## Returning to the IPV example

---

## Load the data


```
ipv <- read_csv("./data/dhs_ipv.csv")  
## on your machine, path is /slides/data/  
head(ipv)
```

```
## # A tibble: 6 x 8
```

```
##      X1 beat_burnfood beat_goesout sec_school no_media countr  
##    <dbl>         <dbl>         <dbl>         <dbl>     <dbl> <chr>  
## 1      1          4.4          18.6          25.2       1.5 Albani  
## 2      4          4.9          19.9          67.7       8.7 Armeni  
## 3      5          2.1          10.3          67.6       2.2 Armeni  
## 4      6          0.3           3.1           46        6.4 Armeni  
## 5      7         12.1          42.5          74.6       7.4 Azerba  
## 6      8          NA           NA            24       41.9 Bangla
```


# Look at bivariate relationships

```
ggplot(ipv, aes(x = sec_school, y = no_media, color = region)) + geom_point()
```

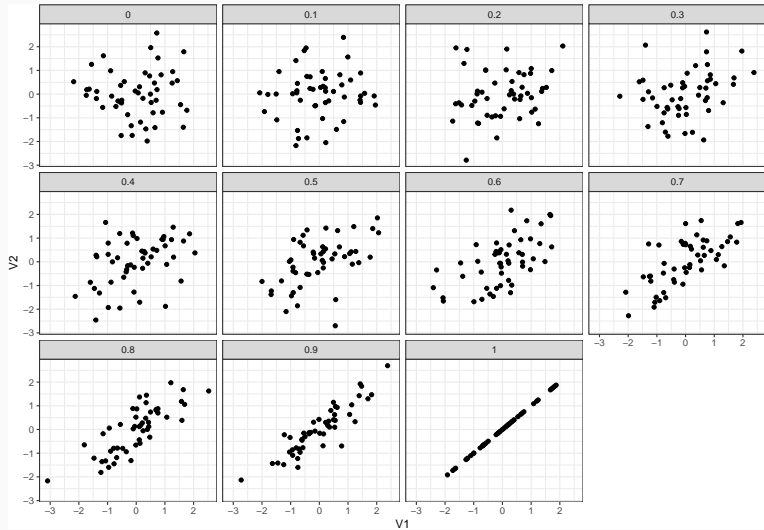


Is there a change in sec\_school by region over time across this sample? Does time matter here?

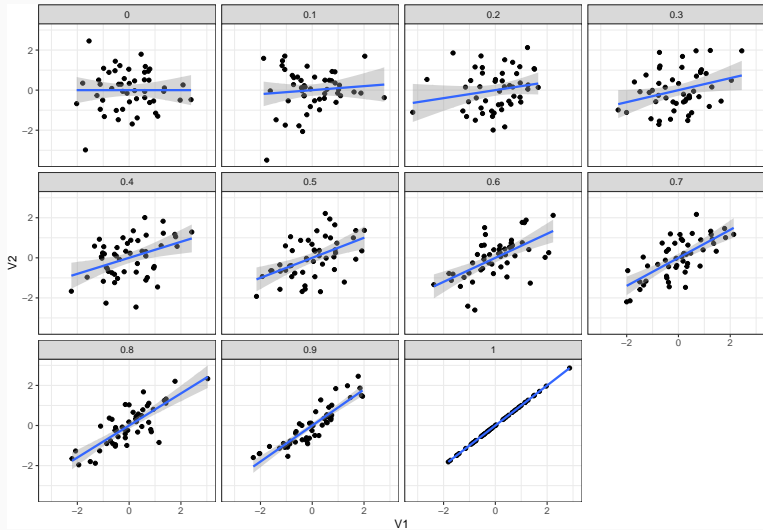
```
ipv_ts <- ipv %>% group_by(region, year) %>% summarise(sec_school = mean(sec_school))  
  
ggplot(ipv_ts, aes(x = year, y = sec_school, color = region)) + geom_line()
```



# Correlation



# Correlation



## Correlation (math time): Z-scores

First, we need the variables to be comparable, so we transform them to be on a standard deviation scale.

A z-score scales a variable measures the number of standard deviations an observation is away from it's mean.

$$\text{z score of } x_i = \frac{x_i - \bar{x}}{S_x}$$

Where  $\bar{x}$  is the mean, and  $S_x$  is the standard deviation of variable  $x$ . Z scores have a mean zero, and a range defined by the range of the data on a standard deviation scale.

For a normally (Gaussian) distributed variable, this will typically range between  $[-3, 3]$

In R, we can transform a numeric into a z-score using `scale()`




## Z-scores in R

```
ipv_scale <- ipv %>% mutate(sec_scale = scale(sec_school)) %>% s  
  sec_scale)  
summary(ipv_scale)
```

##	sec_school	sec_scale.V1
##	Min. : 3.10	Min. : -1.345006
##	1st Qu.: 10.18	1st Qu.: -0.898292
##	Median : 22.40	Median : -0.126408
##	Mean : 24.40	Mean : 0.000000
##	3rd Qu.: 34.90	3rd Qu.: 0.662840
##	Max. : 74.60	Max. : 3.169492
##	NA's : 3	NA's : 3

## Z-scores in R

```
ggplot(ipv_scale, aes(x = sec_school, y = sec_scale)) + geom_point()
```




Correlation measures the degree to which two variables are associated with each other. We often use the letter  $r$  to denote a correlation.

$$r(x, y) = \frac{1}{n} \sum_{i=1}^n \frac{x_i - \bar{x}}{S_x} \times \frac{y_i - \bar{y}}{S_y}$$

Note that this is equal to the average of the product of the *z-scores* of  $x$  and  $y$

In R, you can use `cor()`

## Returning to our example: Are sec\_school and no\_media correlated?



## Obtaining the correlation coefficient

```
cor(ipv$sec_school, ipv$no_media, use = "complete")
```

```
## [1] -0.6077951
```

```
## z score method
```

```
mean(scale(ipv$sec_school) * scale(ipv$no_media), na.rm = TRUE)
```

```
## [1] -0.6084724
```

Data often *cluster* based on unobserved or unobservable characteristics. We can use *classification methods* to try to uncover these latent structures in data.

$k$ -means is a straightforward method we can use to identify  $k$  latent groupings in our data, based on proximity of observations for specified variables.

# The k-means algorithm

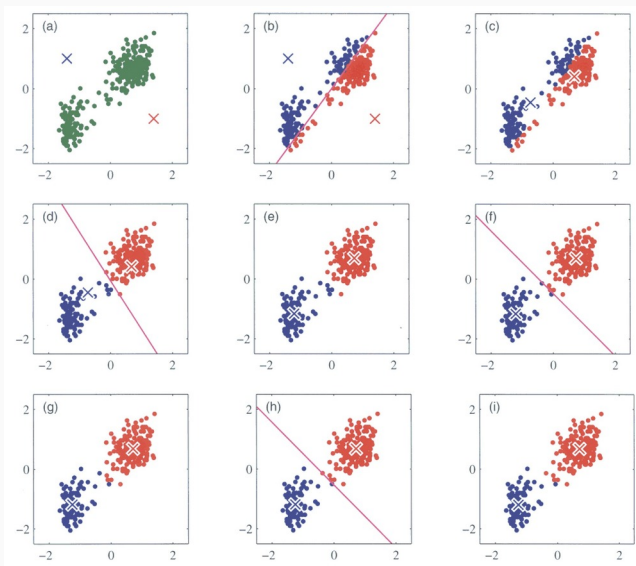
An algorithm is a sequential set of steps used to solve a problem.

A *centroid* is the mean value of a cluster within a group.

1. Choose the initial centroids for each of the  $k$  clusters
2. Assign each observation to the cluster with the nearest centroid
3. Assign a new centroid based on the within-cluster mean for assigned observations
4. Repeat steps 2 and 3 until the cluster assignments no longer change

We arbitrarily choose the number of clusters  $k$ , and R randomly selects starting centroid values for step 1.

# The k-means algorithm





# Implementing k-means for the IPV data

```
ipv_kmeans <- ipv %>% select(sec_school, no_media) %>% mutate(sec_school = scale(sec_school),  
  no_media = scale(no_media)) %>% filter(!(is.na(sec_school)), !(is.na(no_media))) %>%  
  kmeans(centers = 3, nstart = 10)
```

```
ipv_kmeans
```

```
## K-means clustering with 3 clusters of sizes 72, 17, 46  
##  
## Cluster means:  
##   sec_school   no_media  
## 1 -0.6135490  0.7910351  
## 2  1.8803248 -1.1669709  
## 3  0.2071354 -0.7678187  
##  
## Clustering vector:  
##   [1] 3 2 2 2 2 1 1 1 1 1 1 3 3 1 1 1 1 3 1 1 1 1 1 1 1 1 3 3 1 1 1 1 3 3 3 3  
##  [36] 3 3 2 1 1 1 1 2 1 3 3 1 1 2 2 1 3 3 3 3 1 3 3 3 2 2 3 3 2 1 3 1 1 1 1  
##  [71] 1 1 1 3 1 1 1 2 3 1 1 3 3 3 1 1 3 1 1 1 1 1 1 3 3 3 2 3 3 3 3 3 1 1  
## [106] 1 3 3 3 3 3 1 1 2 1 1 1 1 1 1 1 2 1 1 1 1 3 2 2 1 1 1 1 1 1  
##  
## Within cluster sum of squares by cluster:  
## [1] 52.858371  8.657364 24.241269  
## (between_SS / total_SS =  68.3 %)  
##  
## Available components:  
##  
## [1] "cluster"      "centers"      "totss"        "withinss"  
## [5] "tot.withinss" "betweenss"    "size"         "iter"  
## [9] "ifault"
```

## Working with the k-means object

```
str(ipv_kmeans)
```

```
## List of 9
## $ cluster      : int [1:135] 3 2 2 2 2 1 1 1 1 ...
## $ centers       : num [1:3, 1:2] -0.614 1.88 0.207 0.791 -1.167 ...
##  .- attr(*, "dimnames")=List of 2
##  .. ..$ : chr [1:3] "1" "2" "3"
##  .. ..$ : chr [1:2] "sec_school" "no_media"
## $ totss        : num 270
## $ withinss     : num [1:3] 52.86 8.66 24.24
## $ tot.withinss : num 85.8
## $ betweenss    : num 184
## $ size         : int [1:3] 72 17 46
## $ iter         : int 3
## $ ifault       : int 0
## - attr(*, "class")= chr "kmeans"
```

## Pull out what we need from the list


```
ipv_clusters <- ipv %>% filter(!(is.na(sec_school)), !(is.na(no_
  mutate(cluster = factor(ipv_kmeans$cluster))
```

```
library(broom)
centers <- tidy(ipv_kmeans)
centers
```


```
## # A tibble: 3 x 5
##       x1      x2  size withinss cluster
##   <dbl> <dbl> <int>    <dbl> <fct>
## 1 -0.614  0.791    72    52.9    1
## 2  1.88  -1.17    17     8.66    2
## 3  0.207 -0.768    46    24.2    3
```

# Plot it!


```
ggplot(ipv_clusters, aes(x = scale(sec_school), y = scale(no_media), color = cluster)) +  
  geom_point() + geom_point(data = centers, aes(x = x1, y = x2), color = "black",  
    size = 4, shape = 3)
```



What if we thought there were 4 clusters?



What if we thought there were 10 clusters?



- Measurement and design matter!
- Always check your data, and think about how unit and item non-response may inform your conclusions
- Think about desirability and other forms of response bias as you interpret your results
- Design visuals and exploratory analyses to check hypotheses about what's going on in the data
- Think about the structure of your data, use descriptive statistics like correlations to describe relationships
- Think about latent structures in your data to capture clustering

- Complete exercise 3.9.2, use `data(vignettes)`
- You can use `geom_vline()` to add vertical lines
- See <https://jrnold.github.io/qss-tidy/measurement.html#quantile-quantile-plot> for an example of making a quantile-quantile plot with `ggplot()`
- You can do this! Use Slack to ask questions when you are stuck!