

Discovery, 1

Frank Edwards

10/23/2019

Text data: setup

```
# Run me first install.packages(c('tm', 'SnowballC', 'wordcloud', 'modelr'))  
library(tm)  
library(SnowballC)  
library(wordcloud)  
library(modelr)
```

Load the federalist papers

```
DIR_SOURCE <- system.file("extdata/federalist", package = "qss")
corpus.raw <- Corpus(DirSource(directory = DIR_SOURCE, pattern = "fp"))
corpus.raw
```

```
## <<SimpleCorpus>>
## Metadata:  corpus specific: 1, document level (indexed): 0
## Content:   documents: 85
```

Federalist 1

```
content(corpus.raw[[1]])
```

```
## [1] "AFTER an unequivocal experience of the inefficiency of the subsisting \n      federal government"
```

Tidy it up

```
corpus.prep <- tm_map(corpus.raw, content_transformer(tolower))  
corpus.prep <- tm_map(corpus.prep, stripWhitespace)  
corpus.prep <- tm_map(corpus.prep, removePunctuation)  
corpus.prep <- tm_map(corpus.prep, removeNumbers)
```

Federalist 1 (post processing)

```
content(corpus.prep[[1]])
```

```
## [1] "after an unequivocal experience of the inefficiency of the subsisting federal government you are c
```

Remove stopwords

```
head(stopwords("english"))
```

```
## [1] "i"      "me"     "my"     "myself" "we"     "our"
```

```
corpus <- tm_map(corpus.prep, removeWords, stopwords("english"))
```

Federalist 1 (post processing)

```
content(corpus[[1]])
```

```
## [1] " unequivocal experience inefficiency subsisting federal government called upon deliberate
```


Stem words into comparable terms

```
corpus <- tm_map(corpus, stemDocument)
content(corpus[[1]])
```

```
## [1] "unequivoc experi ineffici subsist feder govern call upon deliber new constitut unit state america"
```

Format into document-term matrix

Obtain measures of term-frequency, how often a given stem-term appears in a document or corpus.

```
dtm <- DocumentTermMatrix(corpus)
dtm.mat <- as.matrix(dtm)
dtm.mat[1:5, 1:5]
```

##	Terms					
## Docs	abl	absurd	accid	accord	acknowledg	
## fp01.txt	1	1	1	1		1
## fp02.txt	0	0	0	0		0
## fp03.txt	2	0	0	1		2
## fp04.txt	1	0	0	1		0
## fp05.txt	0	0	0	0		0

Topic discovery is a form of exploratory data analysis that looks at the frequency of terms used in a text or corpus to identify topics discussed in documents.

- Unsupervised learning approaches are agnostic, and look for patterns in the data without identified outcome or predictor variables.
Examples: clustering, topic discovery
- Supervised learning uses theory or hypotheses to explore data looking for particular patterns across outcomes and predictors. Examples: regression

Assumptions: *texts are bags of words*

Topic discovery: wordcloud for Federalist 12

```
wordcloud(colnames(dtm.mat), dtm.mat[12, ], max.words = 20)
```



Topic discovery: wordcloud for Federalist 24

```
wordcloud(colnames(dtm.mat), dtm.mat[24, ], max.words = 20)
```

state
establish
necessnation will
even natur
object two appear
armi peac
must one upon
time without
power garrison
legislatur

The term frequency-inverse document frequency penalizes terms that occur frequently across documents in the corpus, highlighting novel terms in particular texts.

$$\text{tf-idf}(w, d) = \text{tf}(w, d) \times \text{idf}(w)$$

$$\text{idf}(w) = \log \left(\frac{n}{\text{df}(w)} \right)$$

Weighted term frequency

Federalist 12

```
dtm.tfidf <- as.matrix(weightTfIdf(dtm))  
head(sort(dtm.tfidf[12, ], decreasing = TRUE), n = 10)
```

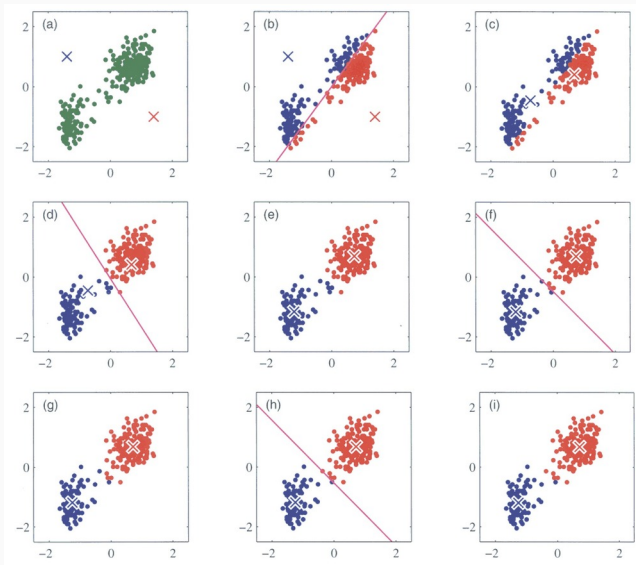
```
##      revenu contraband      patrol      excis      coast      trade  
## 0.01905877 0.01886965 0.01886965 0.01876560 0.01592559 0.01473504  
##      per      tax      cent      gallon  
## 0.01420342 0.01295466 0.01257977 0.01257977
```

Federalist 24

```
dtm.tfidf <- as.matrix(weightTfIdf(dtm))  
head(sort(dtm.tfidf[24, ], decreasing = TRUE), n = 10)
```

```
##      garrison settlement      dockyard      spain      armi      frontier  
## 0.02965511 0.01962294 0.01962294 0.01649040 0.01544256 0.01482756  
##      arsenal      western      post      nearer  
## 0.01308196 0.01306664 0.01236780 0.01166730
```

Remember the k-means algorithm?



We can use k-means to group texts by theme

K-means is an unsupervised approach to uncovering groupings in the data. We can use it to learn and group texts by topic across the corpus.

```
k <- 4 # clusters
# hamilton authored papers (known)
hamilton <- c(1, 6:9, 11:13, 15:17, 21:36, 59:61, 65:85)
dtm.tfidf.hammy <- dtm.tfidf[hamilton, ]
# k-means
km.out <- kmeans(dtm.tfidf.hammy, centers = k)
```

Evaluating clusters

```
for (i in 1:k) {
  print("top words")
  print(head(sort(km.out$centers[i, ], decreasing = TRUE), n = 10))
  print(c("included texts in cluster", i))
  print(rownames(dtm.tfidf.hammy)[km.out$cluster == i])
}

## [1] "top words"
##      pardon      treason      guilt      clemenc      conniv      crime
## 0.04472060 0.02894567 0.02510566 0.02367348 0.02367348 0.01929712
##      impun      plead      sedit      weak
## 0.01788824 0.01673710 0.01492075 0.01470109
## [1] "included texts in cluster" "1"
## [1] "fp74.txt"
## [1] "top words"
##      senat      presid      governor      appoint      nomin      vacanc
## 0.019382349 0.015789668 0.009857989 0.009838966 0.009551661 0.009328505
##      offic      impeach      fill      treati
## 0.007941282 0.006589793 0.006552566 0.006460916
## [1] "included texts in cluster" "2"
## [1] "fp66.txt" "fp67.txt" "fp68.txt" "fp69.txt" "fp75.txt" "fp76.txt"
## [7] "fp77.txt" "fp79.txt"
## [1] "top words"
##      armi      militia      militari      navig      disciplin      war
## 0.011624485 0.011450433 0.008761049 0.005321748 0.004948897 0.004854514
##      peac      northern      frontier confederaci
## 0.004668017 0.004661314 0.004559462 0.004540867
## [1] "included texts in cluster" "3"
## [1] "fp06.txt" "fp08.txt" "fp11.txt" "fp13.txt" "fp24.txt" "fp25.txt"
```

Predicting authorship of federalist papers

66 essays have a known author (Hamilton or Madison). Authorship of the remaining 11 is unknown.

We will use writing style to predict who authored these 11 essays, based on word counts from the known-authorship essays.

```
dtm1 <- as.matrix(DocumentTermMatrix(corpus.prep))  
### count words per 1,000  
tfm <- dtm1/rowSums(dtm1) * 1000  
## subset to words of interest  
words <- c("although", "always", "commonly", "consequently", "considerable",  
           "enough", "there", "upon", "while", "whilst")  
tfm <- tfm[, words]
```

Calculate word usage by Hamilton, Madison

```
madison <- c(10, 14, 37:48, 58)
tfm.ave <- rbind(colSums(tfm[hamilton, ])/length(hamilton), colSums(tfm[madison,
  ])/length(madison))
tfm.ave
```

```
##           although    always  commonly consequently considerable    enough
## [1,] 0.01756975 0.7527744 0.2630876  0.02600857    0.5435127 0.3955031
## [2,] 0.27058809 0.2006710 0.0000000  0.44878468    0.1601669 0.0000000
##           there      upon      while      whilst
## [1,] 4.417750 4.3986828 0.3700484 0.007055719
## [2,] 1.113252 0.2000269 0.0000000 0.380113114
```

Hamilton likes there, upon. Madison likes whilst, consequently. We'll use these style differences to predict who wrote the unknown papers

Predict authorship for unknown papers

Fit a regression model with authorship as the outcome for known papers using distinctive words for each author as predictor. Then predict authorship for unknown papers.

```
author <- rep(NA, nrow(dtm1))  
author[hamilton] <- TRUE  
author[madison] <- FALSE  
author_dat <- data.frame(author = author, tfm) %>% filter(!is.na(author))
```

Fit the model

Fit a model to known authorship papers

```
m0 <- lm(author ~ upon + there + consequently + whilst, data = author_dat)
coef(m0)
```

```
## (Intercept)      upon      there consequently      whilst
##  0.36855800  0.08338945  0.04746947 -0.22006171 -0.32937544
```

Recall that upon, there are Hamilton's words, and consequently, whilst are Madison's. Regression coefficients line up there.

How well does the model fit the data?

- Evaluate total error of the model

- Root Mean Square Error: $\sqrt{\frac{\sum_{i=1}^N (\hat{y}_i - y_i)^2}{N}}$

- Coefficient of determination: $R^2 = 1 - \frac{\sum_{i=1}^n (\hat{y}_i - y_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$

- Evaluate prediction performance

	Positive, obs.	Negative, obs.
Positive, pred.	True positive	False positive
Negative, pred.	False negative	True negative

How well did the model predict authors?

```
author_dat %>% mutate(yhat = fitted(m0) > 0.5) %>% group_by(author) %>% summarise(prop.true.positive = mean(yhat == TRUE, na.rm = TRUE))
```

```
## # A tibble: 2 x 2
##   author prop.true.positive
##   <lgl>         <dbl>
## 1 FALSE         1
## 2 TRUE          1
```

Perfection!

However, we may have *overfit* the data. Fitting your model to a particular dataset is no guarantee that it will predict well *out-of-sample*

We can hold out subsets of the data, re-fit the model, and check classification performance. This gives us a sense of how well our model performs at predicting new cases.

Leave-one-out cross validation: The algorithm

For each row of the data $i = 1 \dots n$

1. Remove i th row of the data
2. Fit the model to the remaining $n - 1$ observations
3. Predict the outcome for the held-out i th observation, calculate prediction error

Then, calculate the mean prediction error across all n observations.

Cross validation using loops

See Arnold for a tidyverse approach using modelr

```
n <- nrow(author_dat)
classify_out <- list()
for (i in 1:n) {
  sub.fit <- lm(author ~ upon + there + consequently + whilst, data = author_dat[-i,
    ])
  classify_out[[i]] <- predict(sub.fit, newdata = author_dat[i, ]) > 0.5
  classify_out[[i]] <- classify_out[[i]] == author_dat$author[[i]]
}
classify_out <- unlist(classify_out)
mean(classify_out)

## [1] 1
```

We still predict perfectly - this is a good model

How about those unidentified cases?

```
author <- rep(NA, nrow(dtm1))
author[hilton] <- TRUE
author[madison] <- FALSE
author_dat <- data.frame(author = author, tfm) %>% filter(!is.na(author))

yhat <- predict(m0, newdata = author_dat)
## Hamilton is >0.5
yhat <- yhat > 0.5
table(yhat)

## yhat
## FALSE TRUE
##      18    1
```

All but one attributed to Madison.

- Next week: network and spatial data
- Homework (two weeks on this one: due 10/30): Choose either text data, network data, or spatial data.
- Text data: 5.5.1, Network data 5.5.2, Spatial data 5.5.3
- Lab: more regression, review HW 4 and prediction from model